

Asset price prediction



using Python, NumPy, Deep Learning and Tensorflow

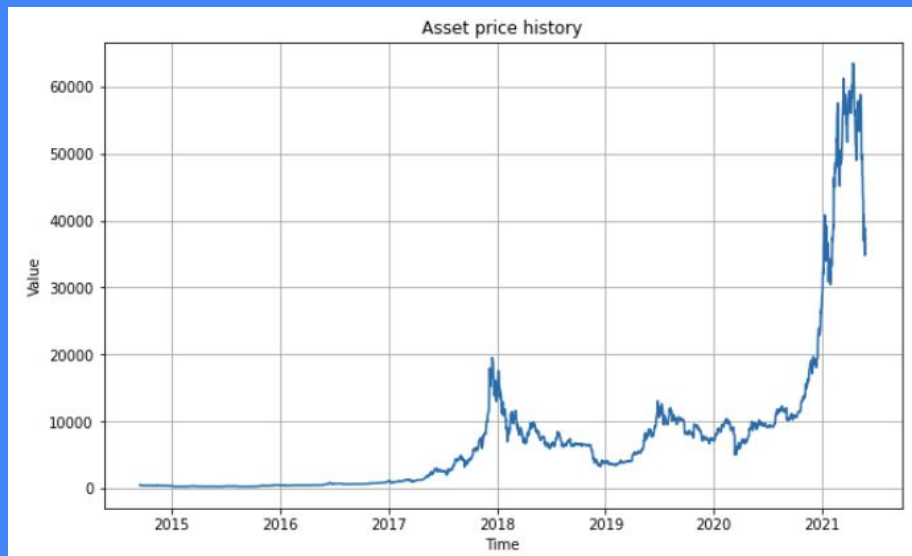
The Plan

- Obtain and prepare the Dataset
- Basic forecasting implementation
 - Naive
 - Moving Average
- Deep Learning implementation
 - Dataset preparation
 - Intro to Deep Learning and Tensorflow
 - Creating the model
 - Training hyperparameters
 - Training the model & forecasting
- Summary & feedback



Disclaimer

This is not a financial advice nor a promotion of trading or any asset. Just sharing my interests, knowledge and experience



Dataset: Timeseries

Exercise: download and explore dataset

- Check you are a member of [Workshops chat](#)
 - Fork & clone [Github repo](#), or use:
 - [Exercises colab notebook](#)
 - [Answers colab notebook](#)
-
- Pick the asset to work with:
 - Heineken, Bitcoin, Dogecoin, Gold, Starbucks, Tesla...
 - Go to finance.yahoo.com and download its trading history csv file
 - Place it under “**sample_data**” directory
 - Explore the dataset
 - ‘Date’ & ‘Close’ columns

Exercise 1: read and prepare data

1	Date	Open	High	Low	Close	Adj Close	Volume
2	2014-09-17	465.864014	468.174011	452.421997	457.334015	457.334015	21056800
3	2014-09-18	456.859985	456.859985	413.104004	424.440002	424.440002	34483200
4	2014-09-19	424.102997	427.834991	384.532013	394.795990	394.795990	37919700
5	2014-09-20	394.673004	423.295990	389.882996	408.903992	408.903992	36863600

Read data from CSV:

```
with open(filepath) as csvfile:

    reader = csv.reader(csvfile, delimiter=',')

    next(reader) # skip the header

    for row in reader: # iterate over rows

        x.append(row[0])

        y.append(row[1])

    return x, y
```

Split into training and validation set:


```
split_time = time_len * 0.8

time_train = time[:split_time]

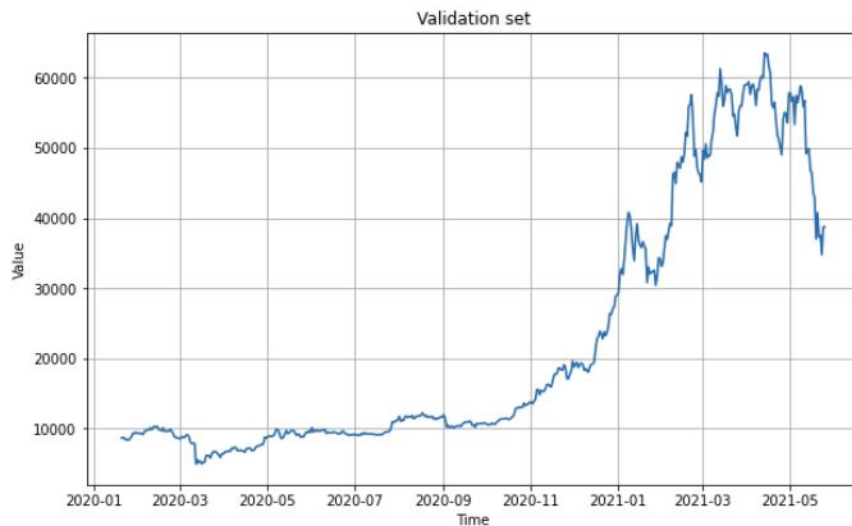
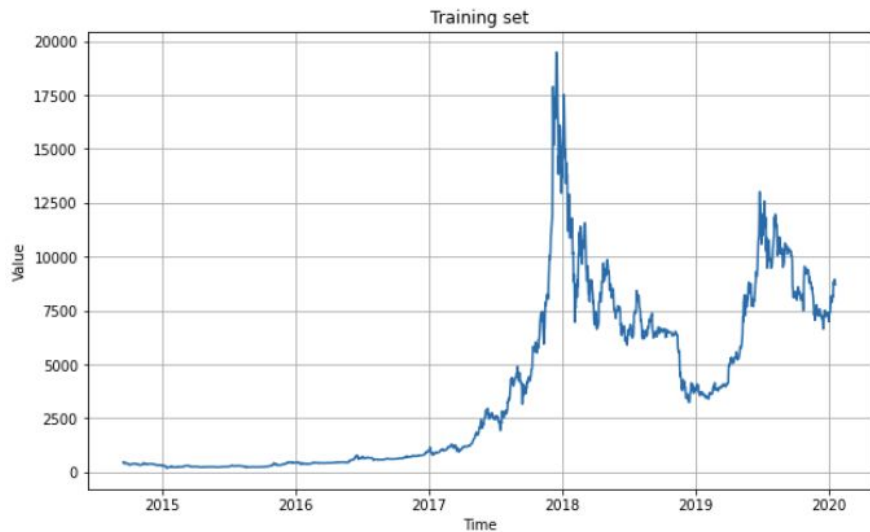
x_train = ...

time_valid = time[split_time:]

x_valid = ...
```



Dataset: plot results





Basic forecasting

Exercise 2: Naive implementation & metrics

- Used as a baseline to compare against other models
- Often is hard to beat
- Basic idea: today's forecast = yesterday's data

```
forecast = series[start - 1:end]
```

- Metrics to evaluate performance:

- MSE - Mean Squared Error

```
keras.metrics.mean_squared_error(validation, forecast)
```

- MAE - Mean Absolute Error


```
keras.metrics.mean_absolute_error(validation, forecast)
```

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

$$\text{MAE} = \frac{1}{n} \sum_{j=1}^n |y_j - \hat{y}_j|$$

Naive implementation: results

MSE:
1567890.9365762253
MAE:
668.2866531270491



Exercise 3: Moving Average

- Simple forecasting method
- Calculates **average** of values over a **fixed period** of time (**averaging window**)
- Usually performs worse than Naive Forecast
- Using window of 30 days:

```
for time in range(series_range - 30):  
    mean = series[time:time + 30].mean()  
    forecast.append(mean)  
  
return forecast
```

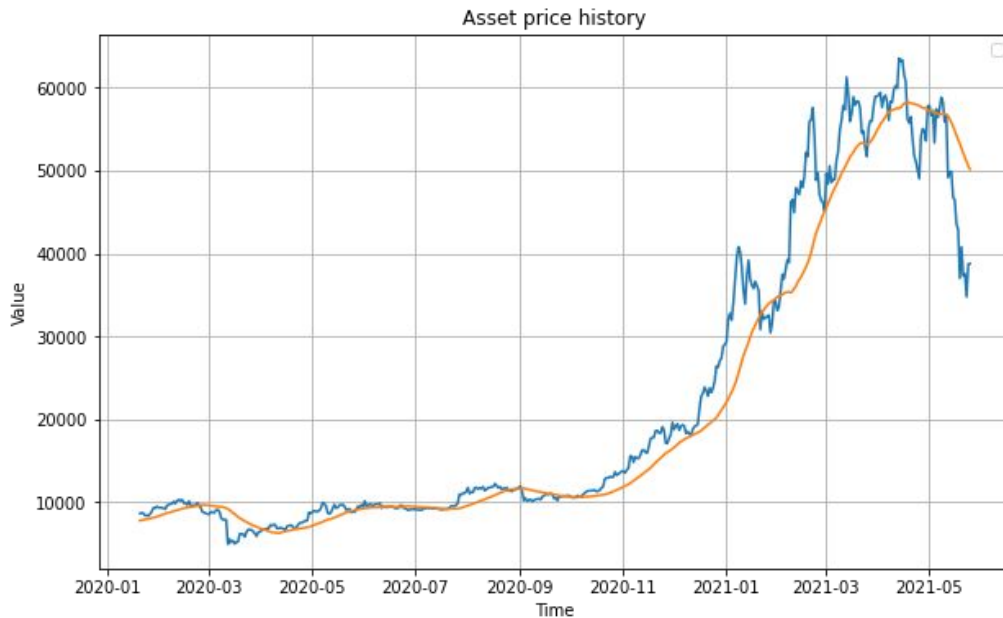
Moving Average: results

Moving Average:

MSE:
17180985.74535265
MAE:
2472.22143907015

Naive:

MSE:
1567890.9365762253
MAE:
668.2866531270491





Deep Learning time!

Exercise 1: Prepare dataset for training


- As before:
 - Use [Github repo](#), or:
 - Use colab [exercises](#) and [answer](#) notebooks
 - Specify and load CSV file
 - Split into training and validation sets

- **Normalize** training and validation sets

```
max_val = np.max(np.abs(my_arr), axis=0)
```

```
my_arr /= max_val
```

- *Split dataset into **smaller windows** of 20 items (input size)




- ***Shuffle** the windows
- *Group in **batches** of 32 windows for parallel training

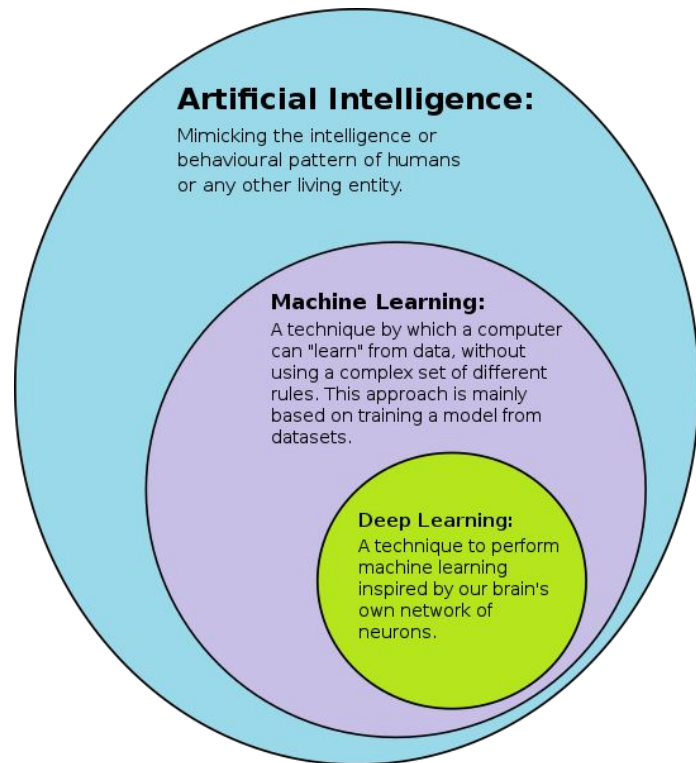
*See `def windowed_dataset` for more details

Intro to Deep Learning and TensorFlow

- Subset of ML based on Artificial Neural Networks



- Applied in many fields:
 - Computer vision
 - Natural language processing
 - Sequence models




Intro to Deep Learning and TensorFlow


- Open source framework to easily create, train & use DL models
- Easy prototyping, deploy & use in Production
- Ecosystem:
 - TensorFlow.js
 - TensorFlow Lite
 - TensorFlow Extended (TFX)
 - TensorFlow Hub
 - TensorFlow Datasets
 - Colab
 - And [many more](#)

Deep Learning model: Perceptron

$$Y = f(WX + b)$$




Deep Learning model: Architecture



Rectified Linear Unit

- Easy to implement and optimize
- Outputs $[0; \infty]$



Exercise 2: Create a DL model

- Do not forget to specify your dataset file
- We will use Tensorflow Sequential API to construct our model
- Example: a simple two-layer network with 10-1 neurons:

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Dense(10, input_shape=[input_size], activation="relu"),  
    tf.keras.layers.Dense(1)  
])  
  
model.summary()
```

Create a DL model: results

Model: "sequential_12"

Layer (type)	Output Shape	Param #
=====		
dense_36 (Dense)	(None, 100)	2100
=====		
dense_37 (Dense)	(None, 10)	1010
=====		
dense_38 (Dense)	(None, 1)	11
=====		

Total params: 3,121

Trainable params: 3,121

Non-trainable params: 0

Training the model

Compile model params:


- Loss function
- Optimizer function
 - Learning rate
 - Momentum

```
model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(lr=1e-6, momentum=0.9))
```

Fit the model on a dataset:

- Dataset to train on
- # of Epochs to train

```
model.fit(dataset, epochs=500, verbose=2)
```



Exercise 3: train the model

Example: compile & train for 50 epochs:

```
model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(lr=1e-5, momentum=0.95))  
  
model.fit(dataset, epochs=50, verbose=2)
```

Training the model: results

Epoch 1/500

61/61 - 0s - loss: 0.0372

Epoch 2/500

61/61 - 0s - loss: 0.0364

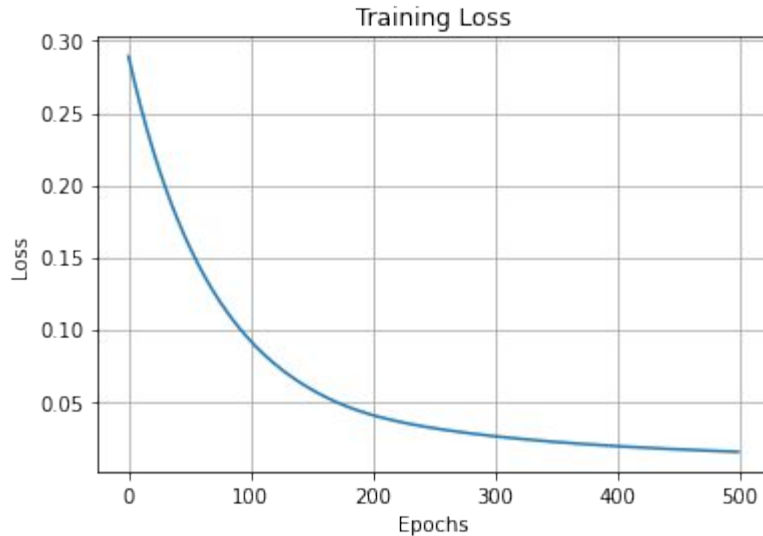
...

Epoch 499/500

61/61 - 0s - loss: 0.0020


Epoch 500/500

61/61 - 0s - loss: 0.0020



Forecasting: results

MSE:
43001720.0
MAE:
4110.766



Summary

- Which model performed better for your data?
- Possible improvements:
 - Play around with **hyperparameters**:
 - # of epochs
 - Adjust learning rate
 - # of neurons/layers
 - Use more **advanced** techniques:
 - Dropout regularization
 - LSTM (Long short-term memory) layers
 - Convolution layers

Feedback

Please [let me know](#) what you think about the topic

Thanks!

	A parrot	Machine learning algorithm
Learns random phrases	✓	✓
Doesn't understand shit about what it learns	✓	✓
Occasionally speaks nonsense	✓	✓
Is a cute birdie parrot	✓	✗